Reference substances: provenance audit

What is documented

The collection metadata records that the target was built with **16 reference substances**. It does not currently provide:

- the exact 16 names;
- the regulatory or scientific source for each name;
- a version or publication date;
- synonyms and spelling variants;
- a justification for the inclusion/exclusion criteria.

Why this matters

Without this information, the target is reproducible only if the original reference list is supplied separately. The current dataset contains `has_risk`, but not the triggering substance(s) for each product.

Responsible decision

This repository does **not** invent or infer the missing list. The project should describe `has_risk` as a supplied, reference-based proxy label. Before production or a stronger academic claim, add a versioned file such as:

Plain Text

```
substance_id,canonical_name,synonyms,source_url,source_date,reason_for_inclusio
```

Then generate `has_risk` and a multi-label target directly from that registry. This will make a deterministic matching baseline possible and auditable.

Deterministic baseline status

A deterministic baseline is the right benchmark for this use case, but it should only be implemented after the exact registry is added. Otherwise, a manually reconstructed list would introduce undocumented assumptions and could produce misleading metrics.